



SLIIT

Discover Your Future

IT2060/IE2061

Operating Systems and System Administration

Lecture 06

Introduction to Deadlock

U. U. Samantha Rajapaksha

M.Sc.in IT, B.Sc.(Engineering) University of Moratuwa

Senior Lecturer SLIIT

Samantha.r@slit.lk



SLIIT
FACULTY OF COMPUTING

DEADLOCKS

When Resources are held by other processes, some processes will wait forever

- Several processes may compete for a finite number of resources, and some of them may wait for the resources forever because the resources are held by other waiting processes → deadlock.
- A set of processes is in a deadlock state if every process in the set is waiting for an event that can be caused only by another process in the set.

Example

- System has two tape drives.
- P1 and P2 each hold one tape drive and each needs another one.

Example

- Semaphores A and B, initialized to 1.

P0

P1

wait (A)

wait (B)

wait (B)

wait (A)

In a situation where, Process 1 cannot release resource 1 for process 2 to acquire because Process 1 cannot be finished without acquiring resource 2 because it has been acquired by process 2. Therefore the lack of resources has been chained and has created a loop

System Model

- Resources are partitioned into several types, each consists of some number of identical *instances*.
 - **Identical:** allocation of *any* instance of the type will satisfy process's request.
 - Resources may be **physical resources** (printers, tape drives, CPU cycles), or **logical resources** (files, semaphores, and monitors).
 - A **pre-emptible** resource is one that can be taken away from a process with no ill effect to the process; e.g., memory.
 - A **non-preemptible** resource is one that cannot be taken away from its user since it will make the user fails; e.g., printers
 - In general, potential deadlocks involve this resource type.
- Each process uses a resource as follows:
 - **Request** the resource; a process must wait if the resource is being used by another process. Request must be sent if a process requires a resource
 - **Use** the resource; e.g., the process can print on the printer.
 - **Release** the resource.

Necessary conditions for deadlock

Four conditions must hold for a deadlock to occur (Coffman et al.):

1. **Mutual exclusion condition.** Only one process at a time can use the resource.

- **or** each resource is either currently assigned to exactly one process or is available.

2. **Hold and wait condition.** A process holding at least one resource is waiting to acquire additional resources held by other processes.

3. **No pre-emption condition.** A resource can be released only voluntarily by the process holding it after that process has completed its task.

4. **Circular wait condition.** There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Note: the four conditions are not completely independent, e.g., the circular-wait condition implies the hold-and-wait condition.

Deadlock Modelling

- Deadlocks can be described more precisely in terms of a directed graph $G(V, E)$
 - called *System resource-allocation graph*
- V is partitioned into two types:
 - Set of processes in the system: $P = \{P_1, P_2, \dots, P_n\}$.
 - Set of all resource types in the system: $R = \{R_1, R_2, \dots, R_n\}$
- *Request edge* – directed edge $P_i \rightarrow R_j$
 - process P_i requests an instance of resource R_j
- *Assignment edge* – directed edge $R_j \rightarrow P_i$
 - an instance of resource R_j has been allocated to process P_i


Model Symbols

- Process: ○

Resource type with 4 instances:




- P_i requests R_j :



R


- P_i uses R_j :



R_j

Example: Deadlock


Printer



Tape drive

- * If the graph contains **no cycles**, no process in the system is deadlocked.
- * If the graph contains **a cycle**, deadlock *may* exist.
 - If each resource type has **one instance**, **cycle means deadlock**.
 - If each resource type has **several instances**, cycle is necessary but **not sufficient condition for deadlock**.

Example: resource allocation graph (with no cycles)



The sets P, R, and E:

$$P = \{P_1, P_2, P_3\}$$

$$R = \{R_1, R_2, R_3, R_4\}$$

$$E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, R_1 \rightarrow P_2, R_2 \rightarrow P_2, R_2 \rightarrow P_1, R_3 \rightarrow P_3\}$$

Resource instances:


- One instance of resource type R_1
- Two instances of resource type R_2
- One instance of resource type R_3
- Three instances of resource type R_4

Process states:

- P_1 is holding an instance of R_2 , and waiting for an instance of R_1
- P_2 is holding an instance of R_1 and R_2 , and is waiting for an instance of R_3
- P_3 is holding an instance of R_3


Example

A cycle and deadlock



A cycle but no deadlock

P_4 can release R_2 which gets allocated to P_3 ; breaking the cycle



Three Methods for handling deadlock

- Use a protocol to ensure that the system will *never* reaches deadlock
 - Using *deadlock prevention* and/or *deadlock avoidance* techniques
- Allow the system to enter a deadlock state and then recover
 - needs *deadlock detection* and *deadlock recovery* algorithms
- Ignore the problem and pretend that deadlocks never occur in the system
 - used by most OS's, including UNIX
 - Also called the **ostrich** algorithm!

Deadlock prevention

- Restrain the ways resource requests can be made
 - Use a set of methods to ensure that **any one** of the four deadlock conditions cannot hold

(1) Deny mutual exclusion

- Not required for sharable resources (e.g., read-only files, cannot be in deadlock)
- Must hold for non-sharable resources (a printer cannot be simultaneously shared by several processes)
- **In general**, it is not possible to prevent deadlock by denying mutual-exclusion condition since some resources are non-sharable

(2) Deny hold and wait

- Must guarantee that whenever a process requests a resource, it does not hold any other resources

Options:

- Each process is granted all resources before it starts
- Allows a process to request resources only when it has none
 - If a process needs more resources, release all resources before requesting new ones

Problem:

- Resource utilisation is low
- Possible starvation.
 - A process that needs popular resources may have to wait indefinitely

Deadlock prevention (cont.)

(3) Prevent no pre-emption (i.e., allow pre-emption)

- When a process holding some resources requests other resource that cannot be immediately allocated, it must release all resources currently being held
 - The pre-empted resources are added to the process's list of requested resources
 - The process is restarted when it regains its old resources and obtains the new one it is requesting

Problem:

- Can be applied easily to resources whose state can be saved easily (e.g., memory), but not so easily for others (e.g., printer)

Deadlock prevention (cont.)

(4) Deny circular wait

- All resource types are ordered, e.g.,
 - $F(\text{card reader}) = 1$ $F(\text{disk drive}) = 5$
 - $F(\text{tape drive}) = 7$ $F(\text{printer}) = 12$
- Each process must request increasing order of resources
- Protocol:
 - Each process requests resources in increasing order
 - Initially a process can request for any R_i
 - After that, it can request R_j only if $F(R_j) > F(R_i)$
- **Problem:** It may be impossible to find a resource ordering that satisfies everyone

Deadlock avoidance

- The system must have some additional *a priori* information about which resources a process will request and use during its lifetime
 - With the additional information, the system can decide for each request whether or not the process should wait
 - The simplest and most useful model requires that each process declare the *maximum* number of resources of each type that it may need
- The deadlock-avoidance algorithm dynamically examines the resource-allocation *state* to ensure that **there can never be a circular-wait condition**
- A resource-allocation *state* is defined by:
 - The number of **available** and **allocated** resources, and
 - The **maximum demands** of the processes


Safe State

- When a process requests an available resource, the system checks if its allocation keeps the system in. *safe state*
- The system is in *safe state* if there exists a *safe sequence* of all processes
- A sequence $\langle P_1, P_2, \dots, P_n \rangle$ is *safe* if, for each P_i , the resources requested by P_i can be allocated from the currently available resources + resources held by all P_j , with $j < i$
 - If P_i 's resource needs are not immediately available, P_i waits until all P_j have finished
 - When all P_j are finished, P_i obtains the needed resources, executes, returns the allocated resources, and terminates
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on

Safe State (cont.)

Basic facts

- If a system is in safe state → no deadlocks
- If a system is in unsafe state → possibility of deadlock
- Avoidance ensures that the system never enters an unsafe state
- A process requesting for a currently available resource may have to wait
 - Thus, resource allocation is lower than without deadlock avoidance algorithm



Example

Consider a system with 12 resources of the same type, and 3 processes with the following resource needs and allocation

	<u>Maximum needs</u>	<u>Allocation</u>	<u>Current need</u>
P_0	10	5	5
P_1	4	2	2
P_2	9	2	7

- At time t_0 , **available resource = 3**, and the system is in safe state
 - There is a safe sequence $\langle P_1, P_0, P_2 \rangle$
- What if at t_1 one more resource is allocated to process P_2 ?
 - The system is in unsafe state
 - Deadlock can occur



Resource-Allocation Graph Algorithm

- *Claim edge* $P_i \rightarrow R_j$ indicates that process P_i may request resource R_j
 - represented by a dashed line
- *Claim edge* converts to *request edge* when a process requests a resource
- When a resource is released by a process, *assignment edge* converts to a *claim edge*
- Resources must be claimed *a priori* in the system
- Need a cycle detection algorithm $\rightarrow O(n^2)$
- This algorithm **can not be used** for system comprising resource types with **multiple instances**

Example

Suppose P_2 requests R_2

Although R_2 is currently free, allocating it to P_2 may lead to unsafe state (a cycle in right figure)



Banker's Algorithm

- The algorithm for a system comprising resource types with **multiple instances**
- Similar to a bank: never allocates its available cash if it can no longer satisfy the needs of all customers
- Each process must *a priori* claim maximum number of instances of each resource type that it may need
- When a process requests a resource:
 - It may have to wait (if resource allocation may lead to unsafe state) until some other process releases enough resources
- When a process gets all its resources:
 - It must return them in a finite amount of time

Banker's Algorithm (cont.)

Algorithm

Let n = number of processes, and m = number of resource types

Data structures:

- *Available*: Vector of length m
 - $available[j] = k$; means k instances of resource type R_j are available
- *Max*: $n \times m$ matrix
 - $Max[i, j] = k$; means process P_i may request at most k instances of resource type R_j .
- *Allocation*: $n \times m$ matrix
 - $Allocation[i, j] = k$; means process P_i is currently allocated k instances of resource type R_j
- *Need*: $n \times m$ matrix
 - $Need[i, j] = k$; means process P_i may need k more instances of resource type R_j to complete its task.
 - $Need[i, j] = max[i, j] - allocation[i, j]$

Implementation of the safety algorithm

// Time complexity = $O(mn^2)$

1. Let *work* and *finish* be vectors of length *m* and *n*, respectively

initialise:

work = *available*

finish [*i*] = *false* for *i* = 1, 2, ..., *n*

// Find an unfinished process *i*; it still needs resources

2. Find a value of *i* such that both:

- *finish*[*i*] = *false*, and
- *need*_{*i*} ≤ *work*
- If no such *i* exists, go to step 4

// process *i* pretends to finish, so it releases its resources i.e., *allocation*_{*i*}

3. *work* = *work* + *allocation*_{*i*}

finish[*i*] = *true*

go to step 2

4. If *finish*[*i*] = *true* for all *i*, the system is in safe state.

Resource-request algorithm for process P_i

$Request_i$ = request vector for process P_i

If $Request_i[j] = k$, then process P_i wants k instances of resource type R_j

1. If $request_i \leq need_i$, go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $request_i \leq available$, go to step 3. Otherwise P_i must wait, since resources are not available.
3. The system pretends to allocate requested resources to P_i by modifying the state as follows:
 - $available = available - request_i$
 - $allocation_i = allocation_i + request_i$
 - $need_i = need_i - request_i$
 - If resulting state is safe, resources are allocated to P_i
 - else P_i must wait, and the old resource-allocation state is restored.

Example of Banker's algorithm

- 5 processes P_0 through P_4 ; 3 resource types A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 :

	<i>Allocation</i>			<i>Max</i>			<i>Available</i>			<i>Need</i>		
	A	B	C	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	5	3	3	3	2	7	4	3
P_1	2	0	0	3	2	2				1	2	2
P_2	3	0	2	9	0	2				6	0	0
P_3	2	1	1	2	2	2				0	1	1
P_4	0	0	2	4	3	3				4	3	1

- ★ The content of matrix *Need* is defined to be *Max* – *Allocation*
- ★ The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies the safety criteria

Example (P_1 requests (1,0,2)):

- Check that $request \leq need$ (that is, $(1, 0, 2) \leq (1, 2, 2) \rightarrow$ true
- Check that $request \leq available$ (that is, $(1, 0, 2) \leq (3, 3, 2) \rightarrow$ true

Before Adjustment

	Allocation		
	A	B	C
P_0	0	1	0
P_1	2	0	0
P_2	3	0	2
P_3	2	1	1
P_4	0	0	2

After Adjustment

	Alloc.			Need			Avail.		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	4	3	2	3	0
P_1	3	0	2	0	2	0			
P_2	3	0	2	6	0	0			
P_3	2	1	1	0	1	1			
P_4	0	0	2	4	3	1			

* $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ or $\langle P_1, P_4, P_3, P_0, P_2 \rangle$ satisfies safety requirement

* Can request for (3, 3, 0) by P_4 be granted? (0, 2, 0) by P_0 ?

Deadlock detection


- If a system does not employ either a deadlock-prevention or a deadlock-avoidance algorithm, then a deadlock situation may occur
- Need a *deadlock detection* algorithm that examines the state of the system to determine whether a deadlock has occurred
- Need a *recovery* algorithm to recover from deadlock

Deadlock detection for single instance of each resource type


- Maintain a *wait-for graph*
 - Nodes are processes
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j
- Periodically invoke an algorithm that searches for a cycle in the graph
 - An algorithm to detect a cycle in a graph requires $O(n^2)$ operations,
 - n is the number of vertices in the graph

Example

Resource Allocation Graph



Wait for Graph



Deadlock recovery

1) Terminate processes

- Kill (abort) all deadlocked processes
- Kill one process at a time until deadlock cycle eliminated
- In which order should we choose process to abort?
 - The process with lowest priority
 - How long the process has computed, and how much longer to completion
 - Resources the process has used
 - Resources the process needs to complete
 - How many processes will need to be terminated
 - Is process interactive or batch?

Problem: what if the process is in the middle of updating a file?

Aborting the process may lead to incorrect file

Deadlock recovery

2) Pre-empt a resource from a process.

- How to select a victim (process) to minimize cost?
- Roll back the process to some safe state and restart from there
 - How do we find a safe state?
 - Easiest way: destroy the process and restart
 - Use checkpoints during execution
- Starvation – same process may always be picked as victim
 - How do we ensure no starvation?
 - Include number of rollbacks in cost factor